Modeling Cognitive Energy from EEG Transitions using Schrodinger Bridge Problem and WGAN-GP

Vaibhav Gollapalli[1]
Texas Academy of Mathematics and Science
University of North Texas, USA
Email: vaibhavgollapalli5@gmail.com

Sriram Sattiraju[1]
Department of Computer Science
University of Texas at Austin, USA
Email: srirams@cs.utexas.edu

Aryan Shah
Texas Academy of Mathematics and Science
University of North Texas, USA
Email: aryanshah.2431@gmail.com

Timothy Fred McMahan
Department of Computer Science
University of North Texas, USA
Email: Fred.McMahan@unt.edu

***Abstract*—Electroencephalography (EEG) provides a non invasive insight into the brain's cognitive and emotional dynamics. However, modeling how these states evolve in real time and quantifying the energy required for such transitions remains a major challenge. The Schrödinger Bridge Problem (SBP) offers a powerful probabilistic framework to model the most efficient evolution between the brain states, interpreted as a measure of cognitive energy cost. In this study, we apply SBP to EEG data collected from STROOP tasks to establish dynamic energy baselines for varying levels of cognitive control and emotional engagement. To address the difficulty of collecting large scale EEG datasets, we also generate synthetic EEG data using a GAN and validate that SBP derived energy metrics are consistent across real and synthetic data. These energy measures can then be used to drive adaptive game environments, enabling real-time adjustments to gameplay difficulty in response to boredom, overstimulation, or other affective states. This approach demonstrates the potential for adaptive interfaces that respond dynamically to a player's cognitive and emotional state.**

***Index Terms*—Electroencephalography (EEG), cognitive dynamics, emotional dynamics, Schrödinger Bridge Problem (SBP), brain state modeling, cognitive energy cost, Stroop task, generative adversarial network (GAN), synthetic EEG data, adaptive gameplay, neuroadaptive interfaces, real-time affective monitoring**

I. Introduction

EEG provides a non-invasive window into neural activity, capturing dynamic oscillatory patterns that reflect cognitive and emotional states in real time. However, quantifying how these neural states evolve, particularly in terms of cognitive energy, remains an unsolved challenge. Cognitive processes such as attention, conflict monitoring, and emotional regulation are inherently dynamic, involving continuous state transitions that cannot be fully represented through static EEG features. Traditional methods for emotion or workload detection rely on static spectral power or connectivity measures, which fail to capture the probabilistic nature of brain-state evolution.

The Schrödinger Bridge Problem (SBP) provides a rigorous mathematical framework to quantify the minimum energy cost required for one probability distribution, representing a brain state, to evolve into another under stochastic dynamics. Originally developed in control theory and optimal transport, SBP has recently been validated in neuroimaging as a novel approach to measure brain state transition costs during cognitive tasks [1]. This method allows for interpreting neural adaptation and mental

effort as an energy optimization problem, aligning well with the concept of cognitive effort in EEG microstate research [2]. In this study we apply SBP to EEG data collected during Stroop tasks, a classical standard of cognitive control and interference resolution. By mapping SBP-derived energy costs to variations in task difficulty, we aim to characterize how neural adaptation unfolds over time. Our goal is to extend these energy based measures into real time adaptive emotion recognition frameworks, suck as gameplay environments that dynamically adjust difficulty or feedback according to the user's affective state. To overcome the limitations of small EEG datasets, especially in emotion and workload research, we also integrate Generative Adversarial Networks (GANs) as a data augmentation mechanism. Prior research demonstrates that GAN-generated EEG can closely replicate real spectral and temporal dynamics [3] [4]. Moreover, recent developments in frequency band conditioned GANs have shown the potential to generate realistic EEG within the alpha, theta, and beta bands, each linked respectively cognitive load, working memory, and mental effort [5]. This work builds upon these advances to establish a foundation for energy aware, adaptive emotion detection using consumer grade EEG devices. By combining the probabilistic structure of SBP with the diversity of the GAN-augmented data, we aim to create a framework capable of continuously tracking cognitive energy and affective state transitions during real world interactive experiences such as gaming.

II. Related Prior Research

A substantial body of work has examined synthetic EEG generation as a solution to limited data availability in neural modeling and Brain computer interfaces research. Early studies demonstrated that GANs can learn complex temporal and spatial EEG distributions and produce synthetic signals suitable for downstream decoding tasks [6]. These foundational efforts established that adversarial learning can model nonstationary neural signals beyond simple noise based augmentation.

GAN-based EEG synthesis was further extended to a variety of paradigms, such as motor imagery, and RSVP tasks. Several studies demonstrated that when synthetic data is added to training pipelines, Wasserstein-based GAN architectures are very good at stabilizing training and maintaining event-related dynamics, which improves classification accuracy [7]. These findings suggest that synthetic EEG can retain cognitively meaningful temporal structure rather than merely approximating spectral statistics.

GAN-based augmentation has also been applied extensively to affective and emotion-related EEG analysis. Several studies report that synthetic EEG improves emotion recognition performance by increasing data diversity and reducing overfitting in deep learning models [8] [9]. Conditional GAN formulations are particularly relevant, as they enable class-consistent signal generation aligned with emotional or cognitive labels, supporting more controlled evaluation of model behavior.

In order to increase the physiological realism of synthetic EEG, more recent research has focused on domain-specific limitations. Alpha, theta, and beta activity can be selectively maintained in artificial signals by conditioning GANs on spectral features linked to cognitive processes using frequency-aware generating techniques. The validity of limited generation procedures is further supported by related clinical approaches that demonstrate conditional Wasserstein GANs' ability to provide diagnostically consistent EEG across illness groups [10] [11]. These studies support the idea that synthetic EEG can encode meaningful neural signatures beyond task-specific decoding.

While pointing out unsolved issues, a number of surveys and toolkits demonstrate the increasing maturity of GAN-based EEG modeling [4] [2]. Reviews highlight that the majority of earlier research assesses synthetic EEG using classifier performance rather than more in-depth dynamical or physiological criteria. Recent frameworks aim to standardize EEG-focused GAN architectures and improve reproducibility across tasks and datasets [12].

Lastly, validation studies of consumer-grade devices have progressively indicated the viability of implementing sophisticated EEG modeling pipelines outside of lab settings. Despite increased

noise and fewer channels, wearable EEG systems can record signals appropriate for cognitive and affective analysis, according to empirical assessments and systematic reviews. [13].This trend motivates investigating whether synthetic data and energy-based metrics remain reliable under realistic acquisition constraints.

Taken together, prior work demonstrates the effectiveness of GANs for EEG data augmentation and the promise of transition-based metrics for modeling cognitive effort, but it leaves open the question of whether synthetic EEG preserves the underlying neural dynamics required for energy-based state transition analysis. Addressing this gap is essential for validating synthetic EEG as a substitute for real recordings in adaptive and real-time neurotechnology systems.

III. Methodology

EEG data was collected during a multi-stage Stroop task designed to elicit progressively increasing levels of cognitive control and interference. Each participant completed four tasks, corresponding to increasing cognitive demand. For each EEG segment, a set of features was extracted to represent instantaneous brain states in a multivariate feature space. These features vectors define empirical probability distributions over neural states for each participant and task condition. To enable meaningful distributional comparison across participants, subject-wise baseline normalization was applied. For each participant, the first Stroop portion was treated as a baseline reference. Given a raw EEG feature vector x, baseline-referenced normalization was performed as:

$$x' = \frac{x - \mu_{\text{baseline}}}{\sigma_{\text{baseline}} + \epsilon} \tag{1}$$

where μ_{baseline} and σ_{baseline} denote the participant-specific mean and standard deviation computed from baseline samples, and ϵ is a small constant added for numerical stability. This normalization removes inter-subject amplitude differences while preserving within-subject deviations associated with task-induced cognitive changes. Since the Schrödinger Bridge Problem (SBP) operates on probability distributions, this step is crucial to ensure that energy estimates reflect cognitive transitions rather than subject specific scaling To further stabilize adversarial training and limit the influence of outliers, normalized features were clipped to a fixed range. All subsequent modeling and SBP analysis were performed in this baseline-referenced feature space.

To address limited EEG sample availability and to model task-conditioned EEG distributions, a conditional Generative Adversarial Network (GAN) was employed. The generator maps samples from a latent noise distribution to EEG feature vectors while conditioning on both participant identity and task portion via learned embeddings. These embeddings are concatenated with the latent input, enforcing subject and task-consistent generation. The generator architecture consists of fully connected layers with residual connections, enabling stable learning in high-dimensional feature space.

The discriminator was implemented as a Wasserstein critic using a packed GAN (PacGAN) formulation. In this approach, multiple samples are concatenated and evaluated jointly by the critic, which reduces mode collapse and improves coverage of multimodal distributions. This design choice is particularly important for EEG data, where realistic modeling requires preserving the overall geometry of the distribution rather than matching marginal statistics alone. Training followed the Wasserstein GAN with gradient penalty (WGAN-GP) framework, which enforces Lipschitz continuity and improves convergence stability. The discriminator loss is given by:

$$\mathcal{L}_D = \mathbb{E}\left[D(\tilde{x})\right] - \mathbb{E}\left[D(x)\right] + \lambda_{\text{GP}}\, \mathbb{E}\left[\left(\|\nabla_{\hat{x}} D(\hat{x})\|_2 - 1\right)^2\right] \tag{2}$$

where $\tilde{x}$ denotes generated samples and $\hat{x}$ represents interpolated samples between real and generated data.

While adversarial training aligns real and synthetic EEG distributions globally, SBP-based energy computation is sensitive to distributional geometry, particularly second-order statistics. To preserve variance structure within each participant and task condition, a condition-wise variance matching regularization term was introduced. For each partic-

ipant–task group present within a training batch, per-feature variances of real and generated samples were computed and penalized when they diverged. The variance matching loss is defined as:

$$\mathcal{L}_{\text{var}} = \frac{1}{\sum_{g\in\mathcal{G}}} \frac{1}{d} \sum_{i=1}^{d} \left(\text{Var}_{\text{real}}^{(g,i)} - \text{Var}_{\text{gen}}^{(g,i)} \right)^2 \quad (3)$$

where $|\mathcal{G}|$ denotes the set of participant–task condition groups and d is the feature dimension. The generator objective combines the adversarial loss and the variance regularization term:

$$\mathcal{L}_G = -\mathbb{E}\left[D(\tilde{x})\right] + \lambda_{\text{var}}\, \mathcal{L}_{\text{var}} \quad (4)$$

This regularization discourages artificial variance collapse and ensure that synthetic EEG distributions preserve covariance structure relevant for energy-based analysis.

Synthetic EEG samples are generated in the baseline-referenced normalized space. For interpretability, samples can be transformed back into original EEG units using the inverse of the subject-specific normalization. However, all distributional comparisons and Schrödinger Bridge computations were performed in the normalized space to ensure consistency between real and synthetic data.

The Schrödinger Bridge Problem was used to quantify the energy cost associated with transitions between EEG-defined brain states. For each participant, empirical distributions P_0 and P_1 were constructed over the EEG feature space for two task portions. SBP computes the minimum stochastic control cost required to transport P_0 to P_1 under entropy-regularized stochastic dynamics. This cost is interpreted as the cognitive energy required for neural adaptation between task conditions.

The SBP formulation and numerical procedure were applied identically to distributions derived from real EEG data and from GAN-generated synthetic data. Agreement between SBP-derived transition energies across real and synthetic datasets serves as validation that the synthetic EEG preserves the distributional structure necessary for meaningful cognitive energy estimation and downstream neuroadaptive applications.

IV. Results

This section presents the experimental results evaluating the fidelity of GAN-generated EEG data and its ability to preserve task-dependent cognitive transport structure as measured by Schrödinger Bridge Process (SBP) energy. We analyze GAN training stability, feature-level distribution alignment, and both group-level and participant-level SBP transport costs, comparing real EEG recordings to GAN-synthesized EEG.

A. GAN Training Stability

To ensure that downstream Schrödinger Bridge Process (SBP) analysis is performed on a reliable generative model, we first evaluate the stability of GAN training. Figure 1 illustrates the evolution of the Wasserstein distance between real and generated EEG distributions across critic update steps, together with a moving average. After an initial transient phase, the Wasserstein distance stabilizes near zero without sustained divergence, indicating convergence of the generator–critic optimization.

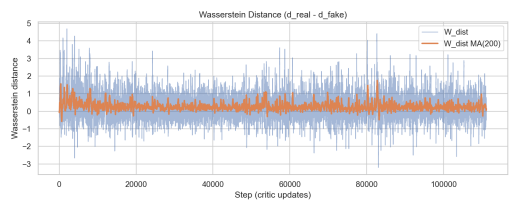

Fig. 1: Wasserstein Distance vs. Critic Updates

Figure 2 shows the corresponding gradient penalty throughout training. The gradient penalty rapidly decreases during early training and remains bounded for the remainder of the optimization, confirming that the Lipschitz constraint is successfully enforced. Together, these results demonstrate stable GAN training without evidence of mode collapse or numerical instability, providing a reliable foundation for subsequent SBP-based transport analysis.

B. Feature-Level Distribution Alignment

We next assess whether GAN-generated EEG preserves the marginal distributions of real EEG

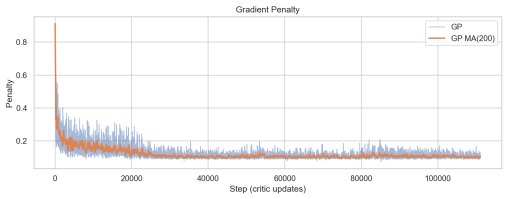

Fig. 2: Gradient Penalty vs. Critic Updates

features. Figure 3 presents overlaid histograms comparing real and synthetic EEG for representative spectral and affective features under baseline z-score normalization, including Theta, Alpha, and Engagement.

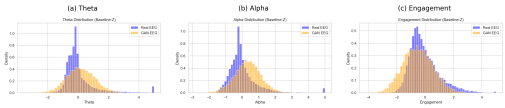

Fig. 3: Overlaid histograms (Real vs. GAN EEG feature distributions)

Across all examined features, GAN-generated EEG closely matches the real EEG distributions in both shape and spread. Real and synthetic data exhibit similar unimodal structure and comparable tail behavior. While small shifts in mean and variance are observed for certain features, these differences are minor relative to the natural inter-participant variability.

Table I quantitatively summarizes the mean and standard deviation for each feature in both real and GAN-generated EEG. The close agreement in first- and second-order statistics indicates that the GAN captures population-level feature distributions without evidence of overfitting or memorization.

TABLE I: Mean and standard deviation of real and GAN-generated EEG features under baseline z-score normalization.

Feature	Real Mean	Real Std	GAN Mean	GAN Std
Theta	0.0797	1.0866	0.1750	0.9550
Alpha	0.0412	0.9363	0.3523	0.8225
Engagement	-0.0627	1.0493	-0.5231	1.0640

C. Schrödinger Bridge Energy on Real EEG

We then apply the Schrödinger Bridge framework to real EEG data to quantify cognitive transport costs between task conditions. SBP energy is computed for transitions from baseline (P1) to subsequent task portions (P2 and P3) for each participant.

Across participants, SBP energy consistently increases from the P1 $\rightarrow$ P2 transition to the P1 $\rightarrow$ P3 transition, reflecting increased cognitive effort with task progression. Although absolute energy values vary between individuals, the direction of change is largely consistent, indicating that SBP energy captures task-dependent cognitive structure in real EEG.

Participant-level SBP energies for real EEG are summarized in Table II.

TABLE II: Participant-level SBP energies for real EEG across task transitions.

Participant	P1 $\rightarrow$ P2 Energy	P1 $\rightarrow$ P3 Energy
Aiden	0.000391	0.000491
Camden	0.000518	0.000549
Emily	0.000725	0.000740
Graham	0.000692	0.000907
Caulin	0.000790	0.001063
Makayla	0.000508	0.000835
Brenda	0.000686	0.000556
Drishti	0.000544	0.000443
Luke	0.000695	0.000694
Maya	0.000858	0.002588

D. Schrödinger Bridge Energy on GAN-Generated EEG

To evaluate whether synthetic EEG preserves transport geometry, we compute SBP energy using GAN-generated EEG under identical conditions. Table III reports SBP energies for GAN EEG across the same task transitions.

At the group level, GAN EEG SBP energies are comparable in magnitude to those computed from real EEG. Mean transport costs for P1 $\rightarrow$ P2 and P1 $\rightarrow$ P3 transitions fall within the same order of magnitude, with overlapping variability. This suggests that GAN-generated EEG preserves the overall geometry of task-dependent cognitive transitions.

TABLE III: Participant-level SBP energies for GAN-generated EEG across task transitions.

Participant	P1 → P2 Energy	P1 → P3 Energy	P1 → P4 Energy
Aiden	0.000845	0.000885	0.000958
Camden	0.000860	0.001127	0.000724
Emily	0.000558	0.000951	0.000777
Graham	0.000753	0.000932	0.001049
Caulin	0.000581	0.000830	0.000666
Makayla	0.000951	0.000948	0.000772
Brenda	0.000693	0.000741	0.000862
Drishti	0.001016	0.000836	0.000827
Luke	0.000740	0.000830	0.001123
Maya	0.000768	0.000779	0.000942

E. Group-Level Comparison of Real and GAN EEG Transport Costs

To directly compare real and synthetic EEG, SBP energies are aggregated across participants and summarized by task transition. Figure 4 presents the mean SBP transport energy with standard deviation for real and GAN EEG for P1 → P2 and P1 → P3 transitions.

For both transitions, GAN EEG exhibits transport energies comparable to real EEG. While GAN EEG shows slightly higher mean SBP energy for the P1 → P2 transition, the difference lies within one standard deviation and does not indicate systematic distortion of transport structure.

These results demonstrate that GAN-generated EEG preserves group-level cognitive transport costs as measured by SBP energy.

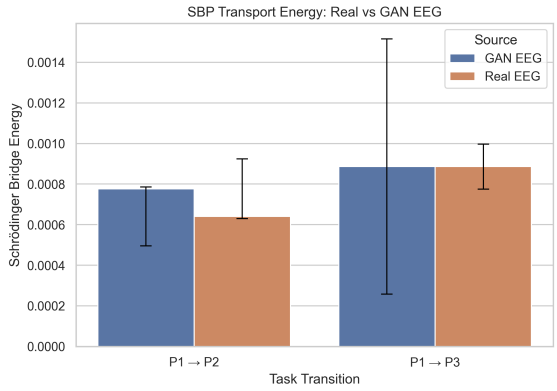

Fig. 4: Bar plot of mean ± standard deviation SBP energy (Real vs. GAN EEG)

F. Preservation of Participant-Level Transport Trends

Beyond group averages, it is critical that synthetic EEG preserves individual-level structure. Figure 5 shows per-participant SBP energy trends for real and GAN EEG across task transitions.

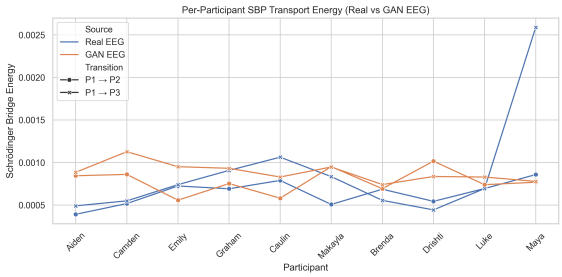

Fig. 5: Bar plot of mean ± standard deviation SBP energy (Real vs. GAN EEG)

For the majority of participants, GAN EEG preserves both the direction and relative ordering of SBP energy changes observed in real EEG. Participants exhibiting larger transport costs in real EEG generally show similarly elevated costs in GAN EEG, even when absolute magnitudes differ.

This consistency indicates that the GAN captures subject-specific cognitive transition geometry rather than collapsing toward population means.

V. Conclusion & Future Works

This work shows that synthetic EEG generated using a conditional WGAN-GP can preserve cognitively meaningful brain state transitions without reproducing EEG signals exactly. By applying the Schrödinger Bridge Problem to both real and generated EEG, we demonstrated that relative transport energy trends across Stroop task conditions are maintained, indicating that the underlying cognitive transport geometry is preserved in synthetic data.

These findings support the use of Schrödinger Bridge transport energy as a principled metric for evaluating synthetic EEG at the distribution level. Future work will extend this approach to continuous task transitions and real-time neuroadaptive systems, enabling dynamic adjustment of interactive environments based on estimated cognitive energy.

REFERENCES

[1] G. Kawakita, S. Kamiya, S. Sasai, J. Kitazono, and M. Oizumi, "Quantifying brain state transition cost via schrödinger bridge," *Netw Neurosci*, vol. 6, no. 1, pp. 118–134, Feb. 2022.

[2] G. Barzon, E. Ambrosini, A. Vallesi, and S. Suweis, "EEG microstate transition cost correlates with task demands," *PLoS Comput Biol*, vol. 20, no. 10, p. e1012521, Oct. 2024.

[3] G. Bouallegue and R. Djemal, "Eeg data augmentation using wasserstein gan," in *2020 20th International Conference on Sciences and Techniques of Automatic Control and Computer Engineering (STA)*, 2020, pp. 40–45.

[4] E. Lashgari, D. Liang, and U. Maoz, "Data augmentation for deep-learning-based electroencephalography," *J Neurosci Methods*, vol. 346, p. 108885, Jul. 2020.

[5] G. Vos, M. Ebrahimpour, L. van Eijk, Z. Sarnyai, and M. R. Azghadi, "A statistical approach for synthetic eeg data generation," 2025. [Online]. Available: https://arxiv.org/abs/2504.16143

[6] F. Fahimi, Z. Zhang, W. B. Goh, K. K. Ang, and C. Guan, "Towards eeg generation using gans for bci applications," in *2019 IEEE EMBS International Conference on Biomedical Health Informatics (BHI)*, 2019, pp. 1–4.

[7] S. Panwar, P. Rad, T.-P. Jung, and Y. Huang, "Modeling eeg data distribution with a wasserstein generative adversarial network to predict rsvp events," *IEEE Transactions on Neural Systems and Rehabilitation Engineering*, vol. 28, no. 8, pp. 1720–1730, 2020.

[8] Y. Luo and B.-L. Lu, "Eeg data augmentation for emotion recognition using a conditional wasserstein gan," in *2018 40th Annual International Conference of the IEEE Engineering in Medicine and Biology Society (EMBC)*, 2018, pp. 2535–2538.

[9] Q. Liu, J. Hao, and Y. Guo, "Eeg data augmentation for emotion recognition with a task-driven gan," *Algorithms*, vol. 16, p. 118, 2023. [Online]. Available: https://api.semanticscholar.org/CorpusID:256965885

[10] A. Mutlu, Şengül Doğan, and T. Tuncer, "Synthetic als-eeg data augmentation for als diagnosis using conditional wgan with weight clipping," 2025. [Online]. Available: https://arxiv.org/abs/2506.16243

[11] F. P. Carrle, Y. Hollenbenders, and A. Reichenbach, "Generation of synthetic EEG data for training algorithms supporting the diagnosis of major depressive disorder," *Front Neurosci*, vol. 17, p. 1219133, Oct. 2023.

[12] J. Song, Q. Zhai, C. Wang, and J. Liu, "EEGGAN-Net: enhancing EEG signal classification through data augmentation," *Front Hum Neurosci*, vol. 18, p. 1430086, Jun. 2024.

[13] H. Rieiro, C. Diaz-Piedra, J. M. Morales, A. Catena, S. Romero, J. Roca-Gonzalez, L. J. Fuentes, and L. L. Di Stasi, "Validation of electroencephalographic recordings obtained with a Consumer-Grade, single dry electrode, Low-Cost device: A comparative study," *Sensors (Basel)*, vol. 19, no. 12, Jun. 2019.